

code

August 17, 2023

0.1 Assignment 1

0.1.1 Name: Ritesah M

0.1.2 Roll Number: 21CS30042

```
[1]: # import all the necessary libraries here
import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
```

```
[2]: df = pd.read_csv('../..../dataset/linear-regression.csv')
print(df.shape)
```

(1599, 12)

```
[3]: df.isna().any()
```

```
[3]: fixed acidity      False
volatile acidity      False
citric acid           False
residual sugar        False
chlorides              False
free sulfur dioxide    False
total sulfur dioxide   False
density               False
pH                    False
sulphates             False
alcohol               False
quality               False
dtype: bool
```

```
[4]: def print_max(dataframe):
    for column in dataframe.columns:
        print(f"Max for {column}: {dataframe[column].max()}")

def print_min(dataframe):
    for column in dataframe.columns:
        print(f"Min for {column}: {dataframe[column].min()}")
```

```
[5]: X=df.drop("quality",axis=1)
      y=df["quality"]
      print(X.shape)
      print(y.shape)
```

```
(1599, 11)
(1599,)
```

```
[6]: print(print_max(df))
      print(print_min(df))
```

```
Max for fixed acidity: 15.9
Max for volatile acidity: 1.58
Max for citric acid: 1.0
Max for residual sugar: 15.5
Max for chlorides: 0.611
Max for free sulfur dioxide: 72.0
Max for total sulfur dioxide: 289.0
Max for density: 1.00369
Max for pH: 4.01
Max for sulphates: 2.0
Max for alcohol: 14.9
Max for quality: 8
None
Min for fixed acidity: 4.6
Min for volatile acidity: 0.12
Min for citric acid: 0.0
Min for residual sugar: 0.9
Min for chlorides: 0.012
Min for free sulfur dioxide: 1.0
Min for total sulfur dioxide: 6.0
Min for density: 0.99007
Min for pH: 2.74
Min for sulphates: 0.33
Min for alcohol: 8.4
Min for quality: 3
None
```

```
[7]: df.columns
```

```
[7]: Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',
          'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density',
          'pH', 'sulphates', 'alcohol', 'quality'],
          dtype='object')
```

```
[8]: cols_to_scale=[]
      for column in df.columns:
```

```

    if column != "quality":
        cols_to_scale.append(column)
print(cols_to_scale)

```

```

['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',
'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density', 'pH',
'sulphates', 'alcohol']

```

```

[9]: from sklearn.model_selection import train_test_split
X_train, X_temp, y_train, y_temp = train_test_split(X, y, test_size=0.5,
↳random_state=42)
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.4,
↳random_state=42)

```

```

[10]: print(X_train.shape)
print(y_train.shape)
print(X_val.shape)
print(y_val.shape)
print(X_test.shape)
print(y_test.shape)

```

```

(799, 11)
(799,)
(480, 11)
(480,)
(320, 11)
(320,)

```

```

[11]: X_train=X_train.to_numpy()
ones_column_train= np.ones((X_train.shape[0],1))
y_train=y_train.to_numpy(dtype="float64")
X_train=np.hstack((ones_column_train,X_train))

X_val=X_val.to_numpy()
ones_column_val= np.ones((X_val.shape[0],1))
y_val=y_val.to_numpy(dtype="float64")
X_val=np.hstack((ones_column_val,X_val))

X_test=X_test.to_numpy()
ones_column_test= np.ones((X_test.shape[0],1))
y_test=y_test.to_numpy(dtype="float64")
X_test=np.hstack((ones_column_test,X_test))

```

1 Analytical Solution

```
[12]: df.columns.shape
```

```
[12]: (12,)
```

```
[13]: df
```

```
[13]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
0	7.4	0.700	0.00	1.9	0.076	
1	7.8	0.880	0.00	2.6	0.098	
2	7.8	0.760	0.04	2.3	0.092	
3	11.2	0.280	0.56	1.9	0.075	
4	7.4	0.700	0.00	1.9	0.076	
...	
1594	6.2	0.600	0.08	2.0	0.090	
1595	5.9	0.550	0.10	2.2	0.062	
1596	6.3	0.510	0.13	2.3	0.076	
1597	5.9	0.645	0.12	2.0	0.075	
1598	6.0	0.310	0.47	3.6	0.067	

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	\
0	11.0	34.0	0.99780	3.51	0.56	
1	25.0	67.0	0.99680	3.20	0.68	
2	15.0	54.0	0.99700	3.26	0.65	
3	17.0	60.0	0.99800	3.16	0.58	
4	11.0	34.0	0.99780	3.51	0.56	
...	
1594	32.0	44.0	0.99490	3.45	0.58	
1595	39.0	51.0	0.99512	3.52	0.76	
1596	29.0	40.0	0.99574	3.42	0.75	
1597	32.0	44.0	0.99547	3.57	0.71	
1598	18.0	42.0	0.99549	3.39	0.66	

	alcohol	quality
0	9.4	5
1	9.8	5
2	9.8	5
3	9.8	6
4	9.4	5
...
1594	10.5	5
1595	11.2	6
1596	11.0	6
1597	10.2	5
1598	11.0	6

[1599 rows x 12 columns]

```
[14]: X_train.shape
```

```
[14]: (799, 12)
```

```
[15]: print(X_train.shape)
      print(y_train)
```

```
(799, 12)
```

```
[5. 5. 7. 6. 6. 5. 6. 5. 6. 5. 5. 5. 6. 5. 7. 7. 5. 5. 5. 5. 5. 5. 6.
 6. 3. 5. 5. 6. 6. 5. 6. 7. 5. 5. 7. 6. 6. 6. 6. 7. 6. 6. 5. 7. 7. 7. 5.
 5. 6. 5. 5. 6. 5. 6. 6. 6. 8. 6. 5. 6. 5. 5. 6. 6. 6. 5. 4. 6. 6. 6. 6.
 6. 6. 5. 5. 8. 4. 5. 5. 6. 6. 6. 7. 5. 5. 6. 6. 6. 6. 6. 6. 5. 5. 6.
 7. 6. 4. 5. 5. 5. 5. 5. 6. 6. 5. 4. 7. 6. 5. 7. 7. 5. 5. 5. 6. 6. 5. 6.
 5. 5. 6. 6. 6. 6. 5. 5. 7. 5. 6. 6. 5. 6. 5. 6. 5. 7. 5. 6. 5. 6. 5. 5.
 6. 6. 5. 5. 6. 6. 7. 6. 5. 5. 7. 6. 5. 5. 6. 7. 6. 6. 6. 7. 5. 6. 7. 6.
 6. 4. 7. 5. 5. 5. 5. 6. 6. 7. 6. 6. 5. 7. 6. 6. 6. 5. 6. 5. 7. 6. 5. 6.
 8. 5. 6. 6. 6. 6. 5. 6. 6. 5. 4. 6. 6. 5. 6. 5. 7. 5. 5. 6. 7. 5. 6. 5.
 6. 5. 6. 5. 7. 5. 6. 6. 6. 7. 6. 8. 6. 5. 6. 5. 6. 4. 5. 5. 6. 7. 5. 5.
 5. 6. 6. 5. 5. 7. 7. 5. 6. 7. 7. 5. 5. 5. 7. 6. 5. 6. 5. 5. 6. 6. 5. 7.
 6. 7. 6. 6. 5. 5. 5. 5. 5. 5. 7. 7. 5. 6. 5. 6. 5. 5. 5. 4. 6. 6. 8. 6.
 6. 5. 5. 5. 5. 6. 6. 5. 5. 5. 5. 5. 5. 7. 5. 5. 5. 5. 4. 5. 6. 6. 5. 5.
 6. 6. 5. 8. 5. 5. 5. 7. 3. 5. 7. 5. 5. 6. 6. 6. 5. 7. 6. 6. 5. 6. 6. 6.
 5. 6. 7. 5. 5. 5. 5. 6. 5. 5. 5. 6. 6. 6. 7. 5. 7. 7. 6. 6. 6. 7. 5. 5.
 6. 6. 5. 5. 4. 6. 6. 6. 6. 5. 5. 6. 4. 6. 5. 7. 6. 6. 6. 7. 4. 6. 5. 5.
 5. 7. 6. 5. 6. 5. 6. 5. 5. 5. 4. 6. 6. 5. 7. 5. 5. 6. 6. 7. 7. 5. 7. 5.
 5. 5. 5. 6. 6. 6. 6. 5. 5. 6. 4. 6. 5. 5. 6. 6. 5. 6. 7. 6. 5. 6. 7. 5.
 6. 7. 5. 6. 5. 6. 6. 6. 6. 5. 6. 5. 5. 6. 5. 5. 6. 5. 5. 5. 6. 4. 6. 5.
 5. 6. 6. 5. 6. 7. 5. 6. 6. 3. 6. 6. 3. 5. 5. 5. 5. 6. 7. 6. 6. 6. 6. 5.
 5. 5. 5. 6. 7. 7. 6. 5. 6. 6. 5. 7. 6. 6. 5. 5. 6. 5. 3. 6. 6. 6. 6. 6.
 6. 6. 5. 5. 6. 5. 7. 5. 6. 7. 5. 5. 7. 6. 6. 6. 5. 7. 6. 6. 5. 4. 6. 5.
 5. 5. 6. 5. 7. 6. 7. 6. 6. 5. 6. 7. 5. 5. 6. 6. 5. 6. 5. 6. 6. 6. 5. 5.
 5. 6. 5. 7. 5. 6. 5. 6. 6. 5. 4. 5. 5. 5. 7. 6. 6. 5. 5. 5. 6. 5. 7. 6.
 6. 5. 5. 5. 5. 5. 5. 6. 5. 6. 5. 6. 6. 5. 5. 5. 5. 5. 6. 7. 5. 6. 6. 6.
 6. 5. 5. 6. 6. 6. 5. 6. 7. 6. 5. 6. 7. 5. 6. 5. 5. 7. 6. 5. 5. 6. 5. 7.
 5. 5. 7. 5. 6. 6. 5. 6. 6. 5. 5. 7. 5. 6. 6. 7. 5. 6. 6. 6. 6. 6. 5. 5.
 5. 7. 7. 5. 6. 6. 6. 5. 4. 5. 6. 5. 6. 6. 5. 5. 5. 6. 6. 7. 5. 7. 7. 5.
 7. 6. 5. 5. 6. 7. 6. 7. 6. 6. 5. 7. 7. 4. 5. 6. 6. 5. 5. 6. 6. 7. 6. 5.
 5. 5. 5. 6. 6. 5. 5. 5. 7. 6. 5. 8. 6. 5. 6. 6. 5. 4. 6. 6. 6. 5. 6. 5.
 5. 6. 6. 6. 6. 7. 5. 5. 6. 7. 3. 6. 6. 6. 4. 6. 6. 4. 6. 6. 5. 6. 4. 6.
 6. 5. 7. 6. 5. 6. 6. 6. 6. 5. 6. 5. 6. 5. 5. 6. 6. 5. 6. 6. 6. 5. 5.
 5. 5. 3. 5. 6. 5. 6. 6. 7. 5. 6. 6. 5. 6. 5. 5. 6. 5. 5. 4. 6. 4. 6. 6.
 6. 5. 6. 6. 5. 7. 6.]
```

```
[16]: Xt_train=np.transpose(X_train)
      print(Xt_train)
```

```
[[ 1.    1.    1.    ...  1.    1.    1.  ]
 [ 8.6   7.8   9.8   ...  7.2   7.9   5.8 ]
 [ 0.52  0.56  0.5    ...  0.62  0.2   0.29]
 ...
 [ 3.2    3.19  3.24 ...  3.51  3.32  3.39]
 [ 0.52  0.93  0.6    ...  0.54  0.8   0.54]
 [ 9.4    9.5   9.7    ...  9.5   11.9  13.5 ]]
```

```
[17]: print(X_train)
```

```
[[ 1.    8.6   0.52 ...  3.2   0.52  9.4 ]
 [ 1.    7.8   0.56 ...  3.19  0.93  9.5 ]
 [ 1.    9.8   0.5    ...  3.24  0.6   9.7 ]
 ...
 [ 1.    7.2   0.62 ...  3.51  0.54  9.5 ]
 [ 1.    7.9   0.2    ...  3.32  0.8   11.9 ]
 [ 1.    5.8   0.29 ...  3.39  0.54  13.5 ]]
```

```
[18]: print(print_max(df))
      print(print_min(df))
```

```
Max for fixed acidity: 15.9
Max for volatile acidity: 1.58
Max for citric acid: 1.0
Max for residual sugar: 15.5
Max for chlorides: 0.611
Max for free sulfur dioxide: 72.0
Max for total sulfur dioxide: 289.0
Max for density: 1.00369
Max for pH: 4.01
Max for sulphates: 2.0
Max for alcohol: 14.9
Max for quality: 8
None
Min for fixed acidity: 4.6
Min for volatile acidity: 0.12
Min for citric acid: 0.0
Min for residual sugar: 0.9
Min for chlorides: 0.012
Min for free sulfur dioxide: 1.0
Min for total sulfur dioxide: 6.0
Min for density: 0.99007
Min for pH: 2.74
Min for sulphates: 0.33
Min for alcohol: 8.4
Min for quality: 3
None
```

```
[19]: XtX_train=np.matmul(Xt_train,X_train)
      XtX_inv_train=np.linalg.inv(XtX_train)
      print(XtX_inv_train)
      print(XtX_inv_train.shape)
```

```
[[ 2.11574955e+03  2.01312462e+00  1.45140448e+00  1.57813370e-01
   9.20446869e-01  3.03217932e+00 -2.62021425e-02  6.21449171e-03
  -2.15934226e+03  1.07472422e+01  2.94881790e+00 -2.00578588e+00]
 [ 2.01312462e+00  3.22026683e-03 -7.49833859e-04 -5.68095055e-03
   8.45858728e-04  1.07236325e-02 -3.95347933e-05  2.13106856e-05
  -2.08656321e+00  1.70945780e-02  2.42460001e-03 -1.89179795e-03]
 [ 1.45140448e+00 -7.49833859e-04  7.61723554e-02  4.86762273e-02
  -1.63186639e-04 -6.86330524e-02  2.10830772e-04 -7.89870899e-05
  -1.50006530e+00 -1.58534290e-04  1.42887160e-02 -5.58737921e-04]
 [ 1.57813370e-01 -5.68095055e-03  4.86762273e-02  9.98141904e-02
  -1.07083492e-03 -7.65867936e-02  2.86912962e-04 -1.31695283e-04
  -1.53067400e-01  5.94837175e-03 -1.30489321e-03 -1.81272642e-03]
 [ 9.20446869e-01  8.45858728e-04 -1.63186639e-04 -1.07083492e-03
   1.06785746e-03 -7.27346859e-05 -2.16787295e-05 -2.26233802e-07
  -9.38606202e-01  4.61728012e-03  1.73878253e-03 -9.90361198e-04]
 [ 3.03217932e+00  1.07236325e-02 -6.86330524e-02 -7.65867936e-02
  -7.27346859e-05  8.64533559e-01 -3.46200057e-04  2.06937659e-04
  -3.41055689e+00  8.24564280e-02 -8.30325562e-02  3.62945602e-03]
 [-2.62021425e-02 -3.95347933e-05  2.10830772e-04  2.86912962e-04
  -2.16787295e-05 -3.46200057e-04  2.25100860e-05 -5.09878419e-06
   2.72403456e-02 -2.85639154e-04 -8.04907241e-05  1.43318450e-05]
 [ 6.21449171e-03  2.13106856e-05 -7.89870899e-05 -1.31695283e-04
  -2.26233802e-07  2.06937659e-04 -5.09878419e-06  2.53519681e-06
  -6.88504971e-03  1.30776574e-04 -5.64659174e-06  6.33502343e-06]
 [-2.15934226e+03 -2.08656321e+00 -1.50006530e+00 -1.53067400e-01
  -9.38606202e-01 -3.41055689e+00  2.72403456e-02 -6.88504971e-03
   2.20568454e+03 -1.14066435e+01 -3.03282330e+00  2.04116816e+00]
 [ 1.07472422e+01  1.70945780e-02 -1.58534290e-04  5.94837175e-03
   4.61728012e-03  8.24564280e-02 -2.85639154e-04  1.30776574e-04
  -1.14066435e+01  1.75223697e-01  1.62852054e-02 -1.27050469e-02]
 [ 2.94881790e+00  2.42460001e-03  1.42887160e-02 -1.30489321e-03
   1.73878253e-03 -8.30325562e-02 -8.04907241e-05 -5.64659174e-06
  -3.03282330e+00  1.62852054e-02  6.02353157e-02 -4.08166660e-03]
 [-2.00578588e+00 -1.89179795e-03 -5.58737921e-04 -1.81272642e-03
  -9.90361198e-04  3.62945602e-03  1.43318450e-05  6.33502343e-06
   2.04116816e+00 -1.27050469e-02 -4.08166660e-03  3.27380605e-03]]
(12, 12)
```

```
[20]: print(XtX_inv_train.shape)
      print(Xt_train.shape)
      print(y_train.shape)
```

(12, 12)

```
(12, 799)
(799,)
```

```
[21]: temp=np.matmul(XtX_inv_train,Xt_train)
      theta=np.matmul(temp,y_train)
      print(theta.shape)
```

```
(12,)
```

```
[22]: print(theta)
```

```
[ 1.03434328e+01  4.21803857e-03 -1.28736172e+00 -2.81156682e-01
  1.59731779e-02 -1.78832680e+00  3.08338676e-03 -3.04438457e-03
 -6.74416241e+00 -1.82876873e-01  6.77925417e-01  2.98166487e-01]
```

```
[23]: y_testpred= np.matmul(X_test,theta)
      print(y_test,y_testpred)
```

```
[5.  6.  6.  6.  6.  4.  6.  6.  6.  7.  5.  5.  5.  5.  6.  6.  6.  5.  5.  5.  6.  6.  5.  6.
  4.  5.  6.  5.  6.  6.  7.  6.  5.  6.  6.  6.  6.  5.  5.  6.  6.  5.  7.  5.  5.  7.  5.  6.
  5.  6.  7.  5.  5.  5.  6.  5.  6.  6.  5.  5.  5.  6.  7.  6.  5.  5.  6.  7.  6.  5.  5.  6.
  5.  5.  5.  7.  6.  7.  4.  5.  7.  5.  6.  5.  5.  6.  5.  5.  5.  5.  7.  5.  6.  4.  6.
  6.  5.  6.  5.  5.  4.  6.  6.  8.  7.  5.  5.  6.  5.  5.  8.  6.  7.  6.  8.  5.  3.  7.  7.
  5.  5.  5.  5.  7.  4.  6.  6.  6.  4.  5.  5.  6.  7.  7.  6.  5.  5.  4.  5.  5.  5.  6.  6.
  5.  6.  5.  6.  6.  5.  6.  6.  7.  5.  5.  5.  5.  5.  3.  5.  7.  5.  6.  5.  7.  6.  6.  5.
  6.  6.  7.  5.  5.  5.  6.  6.  6.  5.  5.  5.  8.  5.  5.  7.  6.  5.  5.  5.  6.  6.  6.  5.
  6.  6.  6.  7.  7.  6.  5.  5.  6.  6.  5.  6.  5.  5.  6.  7.  6.  6.  5.  5.  5.  5.  5.  5.
  6.  7.  6.  6.  5.  5.  5.  6.  7.  6.  5.  5.  5.  6.  7.  6.  5.  6.  5.  6.  7.  5.  5.  6.
  6.  7.  5.  5.  5.  5.  6.  6.  6.  6.  4.  7.  7.  5.  5.  6.  6.  6.  6.  5.  5.  7.  5.  5.
  6.  6.  7.  8.  5.  6.  6.  7.  5.  8.  5.  5.  6.  5.  7.  6.  5.  4.  6.  5.  5.  6.  7.  6.
  5.  6.  7.  6.  6.  5.  7.  6.  6.  5.  5.  6.  4.  6.  6.  6.  5.  5.  6.  3.  5.  5.  7.  6.
  6.  7.  7.  5.  5.  6.  5.  5.] [5.05866571 6.16780205 5.40540303 5.5587164
```

```
6.34309566 5.47243099
```

```
6.23112219 6.22778397 5.22632697 6.56664723 5.47137815 5.04981167
5.08841755 5.34502665 5.66507287 5.85336438 5.13460974 6.03788258
6.01231499 4.73760647 5.20884616 6.12424315 4.54138508 5.3465968
5.93024333 5.29598383 6.18249022 5.2839856 6.27806872 6.48103965
5.85689079 6.18981921 5.24467345 5.15203381 6.25440699 6.47163912
5.43028685 5.29535933 5.4096561 6.32866565 5.74302167 5.67304882
5.87650739 5.06381911 6.33594275 6.62716078 5.19866548 6.0074666
5.57699809 5.51470946 6.27992503 5.08161229 5.38434709 4.9997587
6.07196624 5.21895798 5.27166371 5.0852831 5.53186062 4.98280495
5.85287489 4.88812305 6.13947496 5.94926338 5.08696223 5.4113101
5.60106369 6.52870559 5.25935226 5.32264856 5.66278397 5.61009236
5.9195928 5.3073922 5.66715885 6.1438282 6.04920607 6.08201398
5.10517716 5.26094303 5.84138975 5.38582386 5.7976847 5.9195928
5.32273604 5.60790839 5.36451711 5.65408242 5.41395585 5.18241661
5.08568262 6.23167082 4.98431127 5.63861754 5.15721123 5.96315626
```



```

6.45006723 5.56211563 5.0379615 4.67909326 6.16433077 5.11701904
5.68213639 6.34083887 6.90697326 6.66878237 5.39383094 4.1880468
5.69332989 5.41699169 5.52291283 6.50088403 5.77369332 6.04579926
5.64288512 6.89487312 5.22974473 5.18753644 5.99815584 6.1438282
5.22200601 5.2837158 5.05262825 5.75308989 6.06925539 5.169148
4.95427843 6.16481292 5.11321693 5.1993043 4.91940753 5.14273371
6.43544309 5.57679563 6.32284723 5.28030362 5.51660094 5.3073922
5.36682984 4.84813203 5.05351499 5.23015223 5.73420014 5.91892274
5.14896726 6.06574572 5.18369058 6.19403055 5.16439175 4.94457414
5.9721172 5.534383 5.41418453 5.78939946 4.96553111 4.76184054
5.02498997 5.39205961 4.29031156 5.31762547 5.85616509 5.74071087
6.16481292 5.98294499 6.22206269 6.56580077 6.03805248 5.37558593
6.44054065 5.4910414 6.68734423 5.38015597 6.33379937 6.06708788
6.55385922 5.7886859 4.95537603 5.2279443 5.2572634 5.2217564
6.16343599 5.10186101 4.94754686 6.67385062 5.32206745 5.67735891
5.09802509 5.58827992 5.14779658 5.79139739 6.41373976 5.3094656
5.38159965 5.849234 6.50770047 6.0410544 6.54397238 6.40373348
5.59363992 5.48523671 5.85668851 5.56849689 5.33220126 5.20976546
5.1964433 5.71679951 6.54524835 6.69908776 5.71090726 5.32267735
5.26881523 5.12683181 5.44940684 5.66805004 5.34886877 5.20740738
5.78559111 6.58518276 5.05658731 5.25924647 5.34318731 5.07197065
5.20738301 5.56165485 5.878091 6.05324232 5.07861559 5.1247638
5.49410713 5.94854628 6.47195298 6.59120605 5.19923269 5.4316484
5.81061245 5.12564497 5.9857701 5.02345741 5.11722292 6.25440699
5.39387465 6.18451178 5.33970195 5.33296715 5.28761188 4.76263204
6.29456807 5.70239947 6.02626143 7.03348091 5.49650457 5.77047795
6.51579649 5.73465739 5.67945232 5.20069553 5.39124599 5.2473494
6.67192742 5.39504536 4.93592009 6.68734423 5.14749707 5.25538739
5.48817315 5.60184627 5.85678724 6.52412433 5.16168764 5.31946323
5.73330713 5.96282129 5.25183403 6.57923107 5.70356959 5.54830594
5.12437265 5.08140482 6.00343281 5.68791321 5.4377205 5.16747436
4.95537603 5.21196137 5.10314296 6.02348341 6.29944255 5.70489069
4.95271918 6.57720474 5.61318076 5.61699339 6.23112219 5.40454309
5.89682178 5.42788502 5.77689781 5.43630053 5.98294499 6.34760929
5.28134244 5.87951821 5.11040479 6.1264466 5.79057013 5.10487856
6.32944717 5.28635644 5.25880501 5.16053471 6.13387606 5.2936595
6.05569245 6.13042551 6.0237445 5.40720179 5.22200601 6.33441606
5.40834416 5.14910353]

```

```

[24]: def print_comparision(y_test,y_testpred):
        for i in range(y_test.shape[0]):
            print(f"Actual value:{y_test[i]}, Prediction: {y_testpred[i]}")
            print("\n")

```

```

[25]: def calculate_rmse(y_actual, y_pred):
        squared_errors = (y_actual - y_pred) ** 2
        mean_squared_error = np.mean(squared_errors)

```

```

    rmse = np.sqrt(mean_squared_error)
    return rmse

def calculate_r_squared(y_actual, y_pred):
    total_variance = np.sum((y_actual - np.mean(y_actual)) ** 2)
    explained_variance = np.sum((y_pred - y_actual) ** 2)
    r_squared = 1 - (explained_variance / total_variance)
    return r_squared

```

```

[26]: # from sklearn.metrics import mean_squared_error, r2_score
rmse= calculate_rmse(y_test,y_testpred)
r_squared=calculate_r_squared(y_test,y_testpred)
print(f"mean squared error:{rmse**2}")
print(f"rmse:{rmse}")
print(f"r_squared:{r_squared}")

```

```

mean squared error:0.4431719294227915
rmse:0.6657115962808455
r_squared:0.3967805083955571

```

2 Gradient Descent Solution

Here I am considering theta as w just for the sake of convenience

```

[27]: X_train

```

```

[27]: array([[ 1. ,  8.6 ,  0.52, ...,  3.2 ,  0.52,  9.4 ],
            [ 1. ,  7.8 ,  0.56, ...,  3.19,  0.93,  9.5 ],
            [ 1. ,  9.8 ,  0.5 , ...,  3.24,  0.6 ,  9.7 ],
            ...,
            [ 1. ,  7.2 ,  0.62, ...,  3.51,  0.54,  9.5 ],
            [ 1. ,  7.9 ,  0.2 , ...,  3.32,  0.8 , 11.9 ],
            [ 1. ,  5.8 ,  0.29, ...,  3.39,  0.54, 13.5 ]])

```

```

[28]: X_train.shape

```

```

[28]: (799, 12)

```

```

[29]: y_train.shape

```

```

[29]: (799,)

```

```

[30]: def mean_squared(y_train,y_pred):
        cost=0
        m= y_train.shape[0]
        for i in range(m):
            cost += (y_train[i]-y_pred[i])**2

```

```

error = cost/(2*m)
return error

```

```

[31]: def gradient_descent(X,y,X_val,y_val,w,alpha,train_losses,val_losses):
    m=X.shape[0]
    n=X.shape[1]
    pred=np.dot(X,w)
    error=pred-y
    gradient= np.dot(X.T,error)/m
    w-=alpha*gradient

    val_pred=np.dot(X_val,w)
    # print("the feature constants are",w)
    # print("\n")
    # y_pred= np.matmul(X,w)
    loss=mean_squared(y,pred)
    # y_valpred= np.matmul(X_val,w)
    loss_val= mean_squared(y_val,val_pred)
    train_losses.append(loss)
    val_losses.append(loss_val)
    return w

```

```

[32]: x=np.array([[0,1,2,3,4,5,5,6],
                  [2,4,4,5,6,6,7,4]])
mult=np.array([0,1,2,3,4,5,6,7])
print(x.shape)
print(mult.shape)
print(np.dot(x,mult).shape)
print(x.T)

```

```

(2, 8)
(8,)
(2,)
[[0 2]
 [1 4]
 [2 4]
 [3 5]
 [4 6]
 [5 6]
 [5 7]
 [6 4]]

```

```

[33]: def linear_regression(X,y,X_val,y_val,w,alpha,steps):
    train_losses=[]
    val_losses=[]
    for iter in range(steps):
        w=gradient_descent(X,y,X_val,y_val,w,alpha,train_losses,val_losses)

```

```

y_valpred=np.dot(X_val,w)
return train_losses,val_losses

```

```
[34]: y_val.shape
```

```
[34]: (480,)
```

```

[35]: def normalize_matrix(X):
        columns_to_normalize = range(1, X.shape[1])
        first_column = X[:, 0]
        normalized_columns = (X[:, columns_to_normalize] - X[:,
columns_to_normalize].mean(axis=0)) / X[:, columns_to_normalize].std(axis=0)
        X_scaled = np.column_stack((first_column, normalized_columns))

        print("Original Matrix:")
        print(X)

        print("\nNormalized Matrix:")
        print(X_scaled)
        return X_scaled
X_train_scaled=normalize_matrix(X_train)

```

Original Matrix:

```

[[ 1.    8.6   0.52 ...  3.2   0.52  9.4 ]
 [ 1.    7.8   0.56 ...  3.19  0.93  9.5 ]
 [ 1.    9.8   0.5   ...  3.24  0.6   9.7 ]
 ...
 [ 1.    7.2   0.62 ...  3.51  0.54  9.5 ]
 [ 1.    7.9   0.2   ...  3.32  0.8   11.9 ]
 [ 1.    5.8   0.29 ...  3.39  0.54 13.5 ]]

```

Normalized Matrix:

```

[[ 1.          0.21467856 -0.07076412 ... -0.77358368 -0.81095735
 -0.97734658]
 [ 1.          -0.26952234  0.15919229 ... -0.84003104  1.54211189
 -0.88353352]
 [ 1.          0.94097992 -0.18574232 ... -0.50779423 -0.35182189
 -0.6959074 ]
 ...
 [ 1.          -0.63267302  0.5041269   ...  1.28628454 -0.69617348
 -0.88353352]
 [ 1.          -0.20899723 -1.91041539 ...  0.02378466  0.79601677
  1.36797994]
 [ 1.          -1.48002461 -1.39301347 ...  0.4889162  -0.69617348
  2.8689889 ]]

```

3 $\alpha = 0.1$


```
[36]: w=np.zeros((12,))  
      print(w)
```

```
[0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
```

```
[37]: train_losses,val_losses=linear_regression(X_train_scaled,y_train,X_val,y_val,w,0.  
      ↪1,1000)
```


```
[38]: plt.plot(train_losses)
```

```
[38]: [<matplotlib.lines.Line2D at 0x203dc3c7f10>]
```



```
[39]: plt.plot(val_losses)
```

```
[39]: [<matplotlib.lines.Line2D at 0x203de624640>]
```



4 $\alpha = 0.01$


```
[40]: w=np.zeros((12,))  
      print(w)
```

```
[0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
```

```
[41]: train_losses_2,val_losses_2=linear_regression(X_train_scaled,y_train,X_val,y_val,w,0.  
      ↪01,3000)
```


```
[42]: plt.plot(train_losses_2)
```

```
[42]: [<matplotlib.lines.Line2D at 0x203de47cfa0>]
```



```
[43]: plt.plot(val_losses_2)
```

```
[43]: [<matplotlib.lines.Line2D at 0x203dc402d70>]
```



5 $\alpha = 0.001$


```
[44]: w=np.zeros((12,))  
      print(w)
```

```
[0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
```

```
[45]: train_losses_3,val_losses_3=linear_regression(X_train_scaled,y_train,X_val,y_val,w,0.  
      ↪001,8000)
```


```
[46]: plt.plot(train_losses_3)
```

```
[46]: [<matplotlib.lines.Line2D at 0x203de52beb0>]
```

```
[47]: plt.plot(val_losses_3)
```

```
[47]: [<matplotlib.lines.Line2D at 0x203df7c7070>]
```



```
[48]: X_test.shape
```

```
[48]: (320, 12)
```

```
[49]: X_test_scaled=normalize_matrix(X_test)
```

```
y_testpred=np.dot(X_test_scaled,w)
print(X_test_scaled.shape)
print(w.shape)
y_testpred.shape
```

Original Matrix:

```
[[ 1.    7.9  0.72 ...  3.4   0.53  9.5 ]
 [ 1.   10.6  0.31 ...  3.26  0.86 10.7 ]
 [ 1.    7.3  0.39 ...  3.41  0.54  9.4 ]
 ...
 [ 1.   10.    0.35 ...  3.23  0.52 12.  ]
 [ 1.    9.9  0.35 ...  3.21  0.5   9.5 ]
 [ 1.    7.4  0.66 ...  3.51  0.56  9.4 ]]
```

Normalized Matrix:

```
[[ 1.          -0.25560867  1.02373839 ...  0.5778992  -0.82038121
```

```

-0.80689022]
[ 1.          1.19636957 -1.08398192 ... -0.30752775  1.42863893
 0.28546248]
[ 1.          -0.5782705  -0.67271942 ...  0.64114398 -0.75222908
-0.89791961]
...
[ 1.          0.87370774 -0.87835067 ... -0.4972621  -0.88853333
 1.46884458]
[ 1.          0.81993077 -0.87835067 ... -0.62375167 -1.02483758
-0.80689022]
[ 1.          -0.52449353  0.71529151 ...  1.27359181 -0.61592483
-0.89791961]]
(320, 12)
(12,)

```

[49]: (320,)

[50]: X_test

```

[50]: array([[ 1.   ,  7.9 ,  0.72, ...,  3.4 ,  0.53,  9.5 ],
 [ 1.   , 10.6 ,  0.31, ...,  3.26,  0.86, 10.7 ],
 [ 1.   ,  7.3 ,  0.39, ...,  3.41,  0.54,  9.4 ],
 ...,
 [ 1.   , 10.   ,  0.35, ...,  3.23,  0.52, 12.   ],
 [ 1.   ,  9.9 ,  0.35, ...,  3.21,  0.5 ,  9.5 ],
 [ 1.   ,  7.4 ,  0.66, ...,  3.51,  0.56,  9.4 ]])

```

[51]: `print(f"r_squared: {calculate_r_squared(y_test,y_testpred)}")`

r_squared: 0.3991593477424942

[52]: `print(f"rmse: {calculate_rmse(y_test,y_testpred)}")`

rmse: 0.6643976588767766

[53]: `print(f"mse:{calculate_rmse(y_test,y_testpred)**2}")`

mse:0.4414242491209416