

Finding the right level of abstraction: Program analysis and the Linux kernel


Julia Lawall (Inria/LIP6-Whisper)

March 8, 2016

Our focus

The Linux kernel:

- ▶ Critical software.
 - Used in embedded systems, desktops, servers, etc.
- ▶ Very large.
 - Over 22 000 .c files.
 - Over 13.6 million lines of C code.
 - Increase of 44% since July 2011 (Linux 3.0).
 - Hundreds of contributors.
- ▶ More and less experienced developers.
 - Maintainers, contributors, developers of proprietary drivers



Critical questions

Bugs seem inevitable:

- ▶ How can we find and fix bugs in the code?

Code must continually evolve:

- ▶ How can we improve security, performance, maintainability, etc?

Critical questions

Bugs seem inevitable:

- ▶ How can we find and fix bugs in the code?

Code must continually evolve:

- ▶ How can we improve security, performance, maintainability, etc?

Requires program analysis, program transformation techniques.

- ▶ At the core of programming languages research.

A little history...

Evolution of program analysis research

Starting point:

- ▶ Idealized imperative imperative or functional languages
- ▶ Emphasis on improving precision.
- ▶ Less attention to scalability.

Evolution of program analysis research

Starting point:

- ▶ Idealized imperative imperative or functional languages
- ▶ Emphasis on improving precision.
- ▶ Less attention to scalability.
- ▶ Messy language features not addressed.
- ▶ Relevance of precise analyses to real code not clear.

Program analysis and operating systems

Late 1990s

- ▶ OS community finding that one size does not fit all
 - Need for programmability, and thus safety.
- ▶ Language community saw a need for validation on real applications
 - Operating systems are complex (Linux 1.0: >120 KLOC)
 - OS correctness really matters

Strategy 1: Reimplementation in safer languages

SPIN: Extensible operating system developed in Modula 3
[Bershad et al.: SOSP 1995]

FoxNet: Network protocol stack in SML
[Biagioni, Harper, Lee: LFP 1994, HOSC 2001]

House: Operating system implemented in Haskell
[Hallgren, Jones, Leslie, and Tolmach: ICFP 2005]

Assessment

- ▶ Interesting issues explored.
- ▶ Little practical usage.
- ▶ Legacy incompatible.

Strategy 2: A safer C

CCured: make existing C programs type safe

- ▶ Validatable pointer operations run unchanged.
- ▶ Runtime metadata and checks for dangerous operations.
- ▶ Some valid code is rejected.
- ▶ [Necula, McPeak, Weimer: POPL 2002]

Cyclone: a safe dialect of C

- ▶ Like C, but annotations on dangerous code.
- ▶ [Jim, et al.:USENIX 2002]

Assessment

- ▶ Overheads remain.
- ▶ Library incompatibilities.

Strategy 3: Unsound, incomplete analysis

Observations:

- ▶ The Linux kernel is large, but many modules are self-contained.
 - E.g., flow-sensitive interprocedural pointer analysis perhaps unnecessary.
- ▶ The Linux kernel has to be understood by humans.
 - Open source development model.

Maybe we can do **less** and get **more**:

- ▶ Less precision.
- ▶ Find more real bugs in more large code bases

Metal [Engler et al., OSDI 2002]

Lightweight, programmable bug detection:

- ▶ State machine to describe bug patterns
- ▶ Patterns to match code fragments
- ▶ Applied on control-flow graph

Metal [Engler et al., OSDI 2002]

Lightweight, programmable bug detection:

- ▶ State machine to describe bug patterns
- ▶ Patterns to match code fragments
- ▶ Applied on control-flow graph

Example:

```
state decl any_pointer v;
```

```
start: { kfree(v) } ==> v.freed;
```

```
v.freed:
```

```
  { *v } ==> v.stop, { err("use after free"); }  
| { kfree(v) } ==> v.stop, { err("double free"); }  
;
```

Assessment

- + Lightweight scanning can process a huge code base
- + Hundreds of bugs found, in Linux and BSD
- + More precise version: SLAM (SDV) from Microsoft
- + Shifted attention to what kind of bugs to look for
 - ▶ Motivated protocol mining using machine learning techniques
- State machine notation unstructured
- Finds bugs but doesn't fix them

Our approach: Coccinelle

Coccinelle

Approach:

- ▶ Static analysis to find patterns in C code.
- ▶ Automatic transformation to fix bugs.
- ▶ User scriptable, based on patch notation (**semantic patches**).
- ▶ <http://coccinelle.lip6.fr/>

Coccinelle

Approach:

- ▶ Static analysis to find patterns in C code.
- ▶ Automatic transformation to fix bugs.
- ▶ User scriptable, based on patch notation (**semantic patches**).
- ▶ <http://coccinelle.lip6.fr/>

Goal: Be accessible to C code developers.

Coccinelle

Approach:

- ▶ Static analysis to find patterns in C code.
- ▶ Automatic transformation to fix bugs.
- ▶ User scriptable, based on patch notation (**semantic patches**).
- ▶ <http://coccinelle.lip6.fr/>

Goal: Be accessible to C code developers.

Find once, fix everywhere.

Bug: !x&y

Author: Al Viro <viro@ZenIV.linux.org.uk>

wmi: (!x & y) strikes again

```
diff --git a/drivers/acpi/wmi.c b/drivers/acpi/wmi.c
```

```
@@ -247,7 +247,7 @@
```

```
    block = &wblock->gblock;
```

```
    handle = wblock->handle;
```

```
- if (!block->flags & ACPI_WMI_METHOD)
```

```
+ if (!(block->flags & ACPI_WMI_METHOD))
```

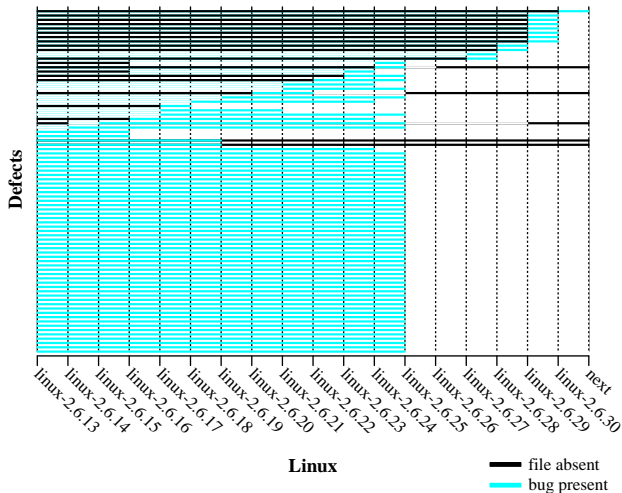
```
    return AE_BAD_DATA;
```

```
if (block->instance_count < instance)
```

Issue

Isolated problems, but these bug types can occur many times

!x&y case:



Bug: !x&y

Author: Al Viro <viro@ZenIV.linux.org.uk>

wmi: (!x & y) strikes again

```
diff --git a/drivers/acpi/wmi.c b/drivers/acpi/wmi.c
```

```
@@ -247,7 +247,7 @@
```

```
    block = &wblock->gblock;
```

```
    handle = wblock->handle;
```

```
- if (!block->flags & ACPI_WMI_METHOD)
```

```
+ if (!(block->flags & ACPI_WMI_METHOD))
```

```
    return AE_BAD_DATA;
```

```
if (block->instance_count < instance)
```

How to automate this change?

- ▶ For any !E & C
 - where E is any expression, and
 - where C is any constant,
- ▶ Add parentheses around E & C

Finding and fixing !x&y bugs using Coccinelle

@@

expression E;

constant C;

@@

- !E & C

+ !(E & C)

- ▶ E is an arbitrary expression.
- ▶ C is an arbitrary constant.

Example

Original code:

```
if (!state->card->
    ac97_status & CENTER_LFE_ON)
    val &= ~DSP_BIND_CENTER_LFE;
```

Semantic patch:

```
@@ expression E; constant C; @@  
- !E & C  
+ !(E & C)
```

Generated code:

```
if (!(state->card->ac97_status & CENTER_LFE_ON))
    val &= ~DSP_BIND_CENTER_LFE;
```

Example

Original code:

```
if (!state->card->
    ac97_status & CENTER_LFE_ON)
    val &= ~DSP_BIND_CENTER_LFE;
```

Semantic patch:

```
@@ expression E; constant C; @@  
- !E & C  
+ !(E & C)
```

Generated code:

```
if (!(state->card->ac97_status & CENTER_LFE_ON))
    val &= ~DSP_BIND_CENTER_LFE;
```

96 instances in Linux from 2.6.13 (August 2005) to
v2.6.28 (December 2008)

Some more Coccinelle features

Dots:

```
a();  
...  
b(x);
```

Nests:

```
if (<+... x == NULL ...+>) S
```

Positions:

```
foo@p(x,y)
```

Some more Coccinelle features

Dots:

```
a();  
... when != x  
b(x);
```

Nests:

```
if (<+... x == NULL ...+>) S
```

Positions:

```
foo@p(x,y)
```

A more complex example

Linux commit 364d5716a:

rtnetlink: ifla_vf_policy: fix misuses of NLA_BINARY

ifla_vf_policy[] is wrong in advertising its individual member types as NLA_BINARY since .type = NLA_BINARY in combination with .len declares the len member as **max** attribute length [0, len].

Our starting point

Linux commit 364d5716a:

rtnetlink: ifla_vf_policy: fix misuses of NLA_BINARY

ifla_vf_policy[] is wrong in advertising its individual member types as NLA_BINARY since .type = NLA_BINARY in combination with .len declares the len member as *max* attribute length [0, len].

```
- [IFLA_VF_MAC]          = { .type = NLA_BINARY,  
-                             .len = sizeof(struct ifla_vf_mac) },  
- [IFLA_VF_VLAN]        = { .type = NLA_BINARY,  
-                             .len = sizeof(struct ifla_vf_vlan) },  
- ...  
+ [IFLA_VF_MAC]          = { .len = sizeof(struct ifla_vf_mac) },  
+ [IFLA_VF_VLAN]        = { .len = sizeof(struct ifla_vf_vlan) },  
+ ...
```

Understanding the problem

Some uses of IFLA_VF_MAC, IFLA_VF_VLAN:

```
struct ifla_vf_mac *ivm = nla_data(tb[IFLA_VF_MAC]);  
struct ifla_vf_vlan *ivv = nla_data(tb[IFLA_VF_VLAN]);
```


Understanding the problem

Some uses of IFLA_VF_MAC, IFLA_VF_VLAN:

```
struct ifla_vf_mac *ivm = nla_data(tb[IFLA_VF_MAC]);  
struct ifla_vf_vlan *ivv = nla_data(tb[IFLA_VF_VLAN]);
```

In a little more detail:

```
struct ifla_vf_mac *ivm = nla_data(tb[IFLA_VF_MAC]);  
  
err = -EOPNOTSUPP;  
if (ops->ndo_set_vf_mac)  
    err = ops->ndo_set_vf_mac(dev, ivm->vf, ivm->mac);
```

Understanding the problem

Some uses of IFLA_VF_MAC, IFLA_VF_VLAN:

```
struct ifla_vf_mac *ivm = nla_data(tb[IFLA_VF_MAC]);  
struct ifla_vf_vlan *ivv = nla_data(tb[IFLA_VF_VLAN]);
```

In a little more detail:

```
struct ifla_vf_mac *ivm = nla_data(tb[IFLA_VF_MAC]);  
  
err = -EOPNOTSUPP;  
if (ops->ndo_set_vf_mac)  
    err = ops->ndo_set_vf_mac(dev, ivm->vf, ivm->mac);
```

Too small ivm may not have vf and mac fields.

Exploring a little more

What are some other usage contexts of `nla_data`?

@@

```
constant c : script:ocaml() { is_nla_binary(c) };  
expression e;
```

@@

```
* nla_data(e[c])
```

Exploring a little more

What are some other usage contexts of `nla_data`?

@@

```
constant c : script:ocaml() { is_nla_binary(c) };  
expression e;
```

@@

```
* nla_data(e[c])
```

Sample result (`drivers/net/macvlan.c`):

```
if (nla_len(tb[IFLA_ADDRESS]) != ETH_ALEN)  
    return -EINVAL;  
if (!is_valid_ether_addr(nla_data(tb[IFLA_ADDRESS])))  
    return -EADDRNOTAVAIL;
```

Assessment

Some observations:

- ▶ `nla_len` obtains the actual data size.
- ▶ Bad sized data should be rejected before accesses.

Assessment

Some observations:

- ▶ `nla_len` obtains the actual data size.
- ▶ Bad sized data should be rejected before accesses.

Semantic patch idea:

- ▶ Find the positions of safe accesses.
- ▶ Report on accesses at other positions.

First attempt

```
@checked1@
constant c;
expression e;
position p;
@@
if (<+... nla_len(e[c]) ...+>) { ... return ...; }
...
nla_data@p(e[c])

@@
constant c : script:ocaml() { is_nla_binary(c) };
expression e;
position p != checked1.p;
@@
* nla_data@p(e[c])
```

Planning ahead

```
@checked2@
```

```
constant c;
```

```
expression l,e,e1;
```

```
position p;
```

```
@@
```

```
l = nla_len(e[c])
```

```
... when != l = e1
```

```
if (<+... l ...+>) { ... return ...; }
```

```
...
```

```
nla_data@p(e[c])
```


More possibilities (simplified)

`nla_data` before test, but test before access:

@checked3@

```
constant c;  
expression e,l,b;  
position p;  
@@
```

```
l = nla_len(e[c])  
b = nla_data@p(e[c])  
if (<+... l ...+>)  
  { ... return ...; }
```

@checked4@

```
constant c;  
expression e,l,b;  
position p;  
@@
```

```
b = nla_data@p(e[c])  
l = nla_len(e[c])  
if (<+... l ...+>)  
  { ... return ...; }
```

Let's try it out!

net/tipc/udp_media.c:

```
if (opts[TIPC_NLA_UDP_LOCAL] && opts[TIPC_NLA_UDP_REMOTE]) {
    sa_local = nla_data(opts[TIPC_NLA_UDP_LOCAL]);
    sa_remote = nla_data(opts[TIPC_NLA_UDP_REMOTE]);
} else {
err:
    pr_err("Invalid UDP bearer configuration");
    return -EINVAL;
}
if ((sa_local->ss_family & sa_remote->ss_family) == AF_INET) {
    ...
}
```

Let's try it out!

net/tipc/udp_media.c:

```
if (opts[TIPC_NLA_UDP_LOCAL] && opts[TIPC_NLA_UDP_REMOTE]) {
    sa_local = nla_data(opts[TIPC_NLA_UDP_LOCAL]);
    sa_remote = nla_data(opts[TIPC_NLA_UDP_REMOTE]);
} else {
err:
    pr_err("Invalid UDP bearer configuration");
    return -EINVAL;
}
if ((sa_local->ss_family & sa_remote->ss_family) == AF_INET) {
    ...
}
```

Let's try it out!

net/tipc/udp_media.c:

```
if (opts[TIPC_NLA_UDP_LOCAL] && opts[TIPC_NLA_UDP_REMOTE]) {
    sa_local = nla_data(opts[TIPC_NLA_UDP_LOCAL]);
    sa_remote = nla_data(opts[TIPC_NLA_UDP_REMOTE]);
} else {
err:
    pr_err("Invalid UDP bearer configuration");
    return -EINVAL;
}
if ((sa_local->ss_family & sa_remote->ss_family) == AF_INET) {
    ...
}
```

Both bugs confirmed.

Some more found code

net/wireless/nl80211.c:

```
if (tb[NL80211_KEY_DATA]) {  
    k->p.key = nla_data(tb[NL80211_KEY_DATA]);  
    k->p.key_len = nla_len(tb[NL80211_KEY_DATA]);  
}
```

Some more found code

net/wireless/nl80211.c:

```
if (tb[NL80211_KEY_DATA]) {  
    k->p.key = nla_data(tb[NL80211_KEY_DATA]);  
    k->p.key_len = nla_len(tb[NL80211_KEY_DATA]);  
}
```

net/netlabel/netlabel_unlabeled.c:

```
ret_val = security_secctx_to_secid(  
    nla_data(info->attrs[NLBL_UNLABEL_A_SECCTX]),  
    nla_len(info->attrs[NLBL_UNLABEL_A_SECCTX]),  
    &secid);
```

Issues

Some uses of `nla_len` and `nla_data` values:

- ▶ Stored in structure fields.
- ▶ Passed to a function.

Issues

Some uses of `nla_len` and `nla_data` values:

- ▶ Stored in structure fields.
- ▶ Passed to a function.

Two options:

- ▶ Collect structure types and field names, and search for uses.
 - Likewise for function names and arguments.

Issues

Some uses of `nla_len` and `nla_data` values:

- ▶ Stored in structure fields.
- ▶ Passed to a function.

Two options:

- ▶ Collect structure types and field names, and search for uses.
 - Likewise for function names and arguments.
- ▶ Assume that if the developer thought of taking the length, he also thought of doing the right thing with it.
 - Less precise, but pragmatic.

Extending the semantic patch: structure fields

```
@semichecked1@
constant c; position p;
expression e,b,bp;
identifier f1,f2; type T;
@@

(
b.f1 = nla_len(e[c]);
...
b.f2 = (T)nla_data@p(e[c]);
|
b.f1 = (T)nla_data@p(e[c]);
...
b.f2 = nla_len(e[c]);
)
```

```
@semichecked2@
constant c; position p;
expression e,b,bp;
identifier f1,f2; type T;
@@

(
b->f1 = nla_len(e[c]);
...
b->f2 = (T)nla_data@p(e[c]);
|
b->f1 = (T)nla_data@p(e[c]);
...
b->f2 = nla_len(e[c]);
)
```

Extending the semantic patch: function arguments

@semichecked3@

```
constant c; position p;  
expression e,e1,e2;  
identifier f; type T;  
@@
```

```
e1 = (T)nla_data@p(e[c])  
...  
e2 = nla_len(e[c])  
...  
(  
f(...,e1,...,e2,...)  
|  
f(...,e2,...,e1,...)  
)
```

@semichecked4@

```
constant c; position p;  
expression e,e1,e2;  
identifier f; type T;  
@@
```

```
e1 = nla_len(e[c])  
...  
e2 = (T)nla_data@p(e[c])  
...  
(  
f(...,e1,...,e2,...)  
|  
f(...,e2,...,e1,...)  
)
```

Also a rule for direct function arguments.

Final bug reporting rule

@checked1@

constant c; expression e; position p;

@@

if (<+... nla_len(e[c]) ...+>) { ... return ...; }

...

nla_data@p(e[c])

[...]

@@

constant c : script:ocaml() { is_nla_binary(c) };

expression e;

position p != { checked1.p, checked2.p, checked3.p, checked4.p,
semichecked1.p, semichecked2.p,
semichecked3.p, semichecked3.p, semichecked5.p }

@@

* nla_data@p(e[c])

Results

- ▶ 15 reports
- ▶ 10 probable real bugs
 - False positives mostly due to separate validation functions
- ▶ Fixes in progress
- ▶ Also possible bugs on `NLA_STRING` and `NLA_NUL_STRING` types

Impact in practice


Methodology

```
git log --grep ... linux
```


- ▶ occinelle
- ▶ semantic patch
- ▶ semantic match
- ▶ SmPL, SMPL, smpl

Use by year


Whole kernel, excluding drivers/staging and scripts:




Use by year: drivers/staging



Use by year: scripts



Who is doing all of this work (up to Sep. 2015)?



Conclusion

Coccinelle: Pragmatic tool for scanning and transforming C code.

- ▶ Improves the reliability of maintenance tasks.
- ▶ Supports both code understanding and bug fixing.
- ▶ Also usable for software metrics.

Conclusion

Coccinelle: Pragmatic tool for scanning and transforming C code.

- ▶ Improves the reliability of maintenance tasks.
- ▶ Supports both code understanding and bug fixing.
- ▶ Also usable for software metrics.

Impact:

- ▶ Around 4000 Coccinelle-based patches accepted into Linux.
- ▶ Over 50 Coccinelle semantic patches available in the Linux source code.
- ▶ Applicable to other C software (wine, qemu, systemd, etc).

Conclusion

Coccinelle: Pragmatic tool for scanning and transforming C code.

- ▶ Improves the reliability of maintenance tasks.
- ▶ Supports both code understanding and bug fixing.
- ▶ Also usable for software metrics.

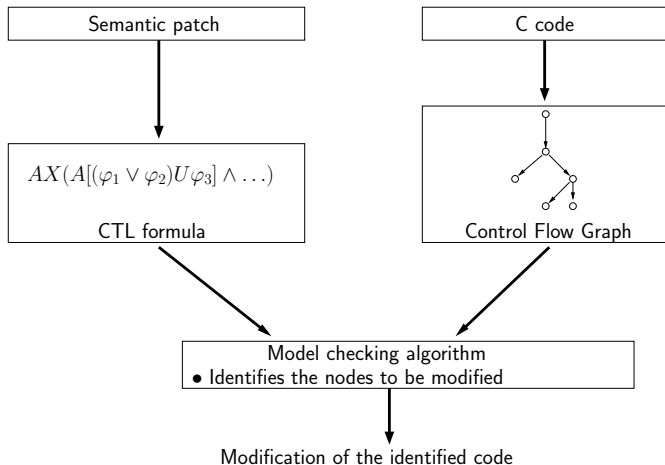
Impact:

- ▶ Around 4000 Coccinelle-based patches accepted into Linux.
- ▶ Over 50 Coccinelle semantic patches available in the Linux source code.
- ▶ Applicable to other C software (wine, qemu, systemd, etc).

<http://coccinelle.lip6.fr/>,
<http://btrlinux.inria.fr/>

How does it work?

Implementation overview



Matching via CTL [POPL'09]

Semantic patch rule \implies CTL formula

Source code function \implies control-flow graph

Provides reasoning about control-flow paths:

- ▶ $a \dots b$ transparently skips over gotos, around loops, etc.
- ▶ Forall (**A**) and exists (**E**) matching available.

CTL for Coccinelle

Extensions:

- ▶ Existentially quantified variables.
- ▶ Witnesses.

CTL for Coccinelle

Extensions:

- ▶ Existentially quantified variables.
- ▶ Witnesses.

Example:

```
@@ expression x; @@  
a();  
...  
b(x);
```

$$a(); \wedge (\mathbf{AX}(\mathbf{A}[\neg(a()); \vee (\exists x, b(x))]) \mathbf{U} (\exists x, b(x))))$$